

# lungcancer

June 28, 2024

```
[ ]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from matplotlib import pyplot as plt
```

```
[ ]: data_true=pd.read_csv("/content/drive/MyDrive/lung cancer/cancer patient data_
↳sets.csv")
df = data_true
```

```
[ ]: df.head()
```

```
[ ]:      index Patient Id  Age  Gender  Air Pollution  Alcohol use  Dust Allergy  \
0         0         P1    33      1             2             4             5
1         1        P10    17      1             3             1             5
2         2       P100    35      1             4             5             6
3         3     P1000    37      1             7             7             7
4         4      P101    46      1             6             8             7
```

```
      OccuPational Hazards  Genetic Risk  chronic Lung Disease  ...  Fatigue  \
0              4              3              2 ...              3
1              3              4              2 ...              1
2              5              5              4 ...              8
3              7              6              7 ...              4
4              7              7              6 ...              3
```

```
      Weight Loss  Shortness of Breath  Wheezing  Swallowing Difficulty  \
0              4              2              2              3
1              3              7              8              6
2              7              9              2              1
3              2              3              1              4
4              2              4              1              4
```

```
      Clubbing of Finger Nails  Frequent Cold  Dry Cough  Snoring  Level
0              1              2              3              4      Low
1              2              1              7              2      Medium
```

2	4	6	7	2	High
3	5	6	7	5	High
4	2	4	2	3	High

[5 rows x 26 columns]

```
[ ]: df.describe()
```

```
[ ]:
```

	index	Age	Gender	Air Pollution	Alcohol use \
count	1000.000000	1000.000000	1000.000000	1000.0000	1000.000000
mean	499.500000	37.174000	1.402000	3.8400	4.563000
std	288.819436	12.005493	0.490547	2.0304	2.620477
min	0.000000	14.000000	1.000000	1.0000	1.000000
25%	249.750000	27.750000	1.000000	2.0000	2.000000
50%	499.500000	36.000000	1.000000	3.0000	5.000000
75%	749.250000	45.000000	2.000000	6.0000	7.000000
max	999.000000	73.000000	2.000000	8.0000	8.000000

	Dust Allergy	OccuPational Hazards	Genetic Risk	chronic Lung Disease \
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	5.165000	4.840000	4.580000	4.380000
std	1.980833	2.107805	2.126999	1.848518
min	1.000000	1.000000	1.000000	1.000000
25%	4.000000	3.000000	2.000000	3.000000
50%	6.000000	5.000000	5.000000	4.000000
75%	7.000000	7.000000	7.000000	6.000000
max	8.000000	8.000000	7.000000	7.000000

	Balanced Diet ...	Coughing of Blood	Fatigue	Weight Loss \
count	1000.000000 ...	1000.000000	1000.000000	1000.000000
mean	4.491000 ...	4.859000	3.856000	3.855000
std	2.135528 ...	2.427965	2.244616	2.206546
min	1.000000 ...	1.000000	1.000000	1.000000
25%	2.000000 ...	3.000000	2.000000	2.000000
50%	4.000000 ...	4.000000	3.000000	3.000000
75%	7.000000 ...	7.000000	5.000000	6.000000
max	7.000000 ...	9.000000	9.000000	8.000000

	Shortness of Breath	Wheezing	Swallowing Difficulty \
count	1000.000000	1000.000000	1000.000000
mean	4.240000	3.777000	3.746000
std	2.285087	2.041921	2.270383
min	1.000000	1.000000	1.000000
25%	2.000000	2.000000	2.000000
50%	4.000000	4.000000	4.000000
75%	6.000000	5.000000	5.000000
max	9.000000	8.000000	8.000000

	Clubbing of Finger Nails	Frequent Cold	Dry Cough	Snoring
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	3.923000	3.536000	3.853000	2.926000
std	2.388048	1.832502	2.039007	1.474686
min	1.000000	1.000000	1.000000	1.000000
25%	2.000000	2.000000	2.000000	2.000000
50%	4.000000	3.000000	4.000000	3.000000
75%	5.000000	5.000000	6.000000	4.000000
max	9.000000	7.000000	7.000000	7.000000

[8 rows x 24 columns]

```
[ ]: df.isnull().sum()
```

```
[ ]: index          0
      Patient Id    0
      Age           0
      Gender        0
      Air Pollution  0
      Alcohol use    0
      Dust Allergy   0
      OccuPational Hazards  0
      Genetic Risk   0
      chronic Lung Disease  0
      Balanced Diet  0
      Obesity        0
      Smoking        0
      Passive Smoker  0
      Chest Pain     0
      Coughing of Blood  0
      Fatigue        0
      Weight Loss    0
      Shortness of Breath  0
      Wheezing       0
      Swallowing Difficulty  0
      Clubbing of Finger Nails  0
      Frequent Cold  0
      Dry Cough      0
      Snoring        0
      Level          0
      dtype: int64
```

```
[ ]:
```

```
data = data_true.drop(["index", "Patient Id", "Age", "Gender", "Level", "Dust",
↳Allergy", "chronic Lung Disease", "Balanced Diet", "Obesity", "Chest",
↳Pain", "Coughing of Blood", "Fatigue", "Weight Loss", "Snoring", "Dry",
↳Cough", "Frequent Cold", "Clubbing of Finger Nails", "Swallowing",
↳Difficulty", "Wheezing", "Shortness of Breath", "Passive",
↳Smoker", "Smoking", "OccuPational Hazards"], axis=1)
data.head(10)
```

```
[ ]:   Air Pollution  Alcohol use  Genetic Risk
0           2           4           3
1           3           1           4
2           4           5           5
3           7           7           6
4           6           8           7
5           4           5           5
6           2           4           3
7           3           1           2
8           4           5           6
9           2           3           4
```

```
[ ]: x = df[['Air Pollution', 'Alcohol use']]
y = df['Genetic Risk']
```

```
[ ]: k = 3
knn = KNeighborsClassifier(n_neighbors=k)
knn.fit(x,y)
```

```
[ ]: KNeighborsClassifier(n_neighbors=3)
```

```
[ ]: # New data point to predict on
new_data = np.array([[6,8]])
prediction = knn.predict(new_data)

# Check the prediction and print the result
if prediction[0] == 0: # Access the prediction string from the array
    print("low Risk")
else:
    print("high Risk")
```

high Risk

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does
not have valid feature names, but KNeighborsClassifier was fitted with feature
names
  warnings.warn(
```

```
[ ]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
[ ]: from sklearn.linear_model import LinearRegression
LR = LinearRegression()
```

```
[ ]: from sklearn.model_selection import train_test_split as ttp
from sklearn.metrics import classification_report
```

```
[ ]: # Assuming 'y' is your target variable and it contains strings like 'yes' and 'no'
# Convert 'yes' to 1 and 'no' to 0
y_numeric = y.replace({'yes': 1, 'no': 0})

# Now fit the model with the numeric target variable
LR.fit(x, y_numeric)
```

```
[ ]: LinearRegression()
```

```
[ ]: new_data = np.array([[6, 35]])
prediction = LR.predict(new_data)[0]
# Check the prediction and print the result
if prediction == 0: # Removed indexing as 'prediction' is a scalar
    print("low Risk")
else:
    print("high Risk")
```

high Risk

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names  
warnings.warn(

```
[ ]: # Install necessary libraries if not already installed
!pip install -q pandas scikit-learn matplotlib

# Import necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
```

```
[ ]: # Load data (replace 'lung_cancer_data.csv' with your actual data file)
data = pd.read_csv('/content/drive/MyDrive/lung cancer/cancer patient data sets.
↳csv')

# Display first few rows to understand the data
print(data.head())

# Define independent variables (features) and dependent variable (target)
X = data[['Air Pollution', 'Alcohol use']].values # Features
y = data['Genetic Risk'].values # Target variable

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,↳
↳random_state=42)

# Standardize features (optional but recommended for logistic regression)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

	index	Patient Id	Age	Gender	Air Pollution	Alcohol use	Dust Allergy	\
0	0	P1	33	1	2	4	5	
1	1	P10	17	1	3	1	5	
2	2	P100	35	1	4	5	6	
3	3	P1000	37	1	7	7	7	
4	4	P101	46	1	6	8	7	

	OccuPational Hazards	Genetic Risk	chronic Lung Disease	...	Fatigue	\
0	4	3	2	...	3	
1	3	4	2	...	1	
2	5	5	4	...	8	
3	7	6	7	...	4	
4	7	7	6	...	3	

	Weight Loss	Shortness of Breath	Wheezing	Swallowing Difficulty	\
0	4	2	2	3	
1	3	7	8	6	
2	7	9	2	1	
3	2	3	1	4	
4	2	4	1	4	

	Clubbing of Finger Nails	Frequent Cold	Dry Cough	Snoring	Level
0	1	2	3	4	Low
1	2	1	7	2	Medium
2	4	6	7	2	High
3	5	6	7	5	High
4	2	4	2	3	High

[5 rows x 26 columns]

```
[ ]: # Initialize logistic regression model
model = LogisticRegression()

# Train the model
model.fit(X_train, y_train)
```

```
[ ]: LogisticRegression()
```

```
[ ]: # Predict on test data
y_pred = model.predict(X_test)

# Print classification report
print(classification_report(y_test, y_pred))

# Plot confusion matrix
conf_mat = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
plt.imshow(conf_mat, cmap=plt.cm.Blues)
plt.title('Confusion Matrix')
plt.colorbar()
plt.xticks([0, 1], ['No Cancer', 'Cancer'])
plt.yticks([0, 1], ['No Cancer', 'Cancer'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/\_classification.py:1344:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero\_division` parameter to  
control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/\_classification.py:1344:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero\_division` parameter to  
control this behavior.


```
_warn_prf(average, modifier, msg_start, len(result))
```

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/\_classification.py:1344:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero\_division` parameter to  
control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

	precision	recall	f1-score	support
1	0.00	0.00	0.00	5

2	0.73	0.67	0.70	49
3	0.34	0.52	0.41	25
4	0.00	0.00	0.00	4
5	1.00	0.93	0.97	15
6	0.00	0.00	0.00	20
7	0.78	0.98	0.86	82
accuracy			0.70	200
macro avg	0.41	0.44	0.42	200
weighted avg	0.62	0.70	0.65	200



```
[ ]: # Install necessary libraries if not already installed
!pip install -q pandas scikit-learn matplotlib

# Import necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, plot_tree
```



```
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
```

```
[ ]: # Load data (replace 'lung_cancer_data.csv' with your actual data file)
data = pd.read_csv('/content/drive/MyDrive/lung cancer/cancer patient data sets.
↳csv')

# Display first few rows to understand the data
print(data.head())

# Define independent variables (features) and dependent variable (target)
X = data[['Air Pollution', 'Alcohol use']].values # Features
y = data['Genetic Risk'].values # Target variable

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,↳
↳random_state=42)
```

	index	Patient Id	Age	Gender	Air Pollution	Alcohol use	Dust Allergy	\
0	0	P1	33	1	2	4	5	
1	1	P10	17	1	3	1	5	
2	2	P100	35	1	4	5	6	
3	3	P1000	37	1	7	7	7	
4	4	P101	46	1	6	8	7	

		OccuPational Hazards	Genetic Risk	chronic Lung Disease	...	Fatigue	\
0		4	3	2	...	3	
1		3	4	2	...	1	
2		5	5	4	...	8	
3		7	6	7	...	4	
4		7	7	6	...	3	

		Weight Loss	Shortness of Breath	Wheezing	Swallowing Difficulty	\
0		4	2	2	3	
1		3	7	8	6	
2		7	9	2	1	
3		2	3	1	4	
4		2	4	1	4	

		Clubbing of Finger Nails	Frequent Cold	Dry Cough	Snoring	Level
0		1	2	3	4	Low
1		2	1	7	2	Medium
2		4	6	7	2	High
3		5	6	7	5	High
4		2	4	2	3	High

[5 rows x 26 columns]

```
[ ]: # Initialize decision tree classifier
clf = DecisionTreeClassifier(random_state=42)

# Train the model
clf.fit(X_train, y_train)
```

```
[ ]: DecisionTreeClassifier(random_state=42)
```

```
[ ]: # Predict on test data
y_pred = clf.predict(X_test)

# Print classification report
print(classification_report(y_test, y_pred))

# Plot confusion matrix
conf_mat = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
plt.imshow(conf_mat, cmap=plt.cm.Blues)
plt.title('Confusion Matrix')
plt.colorbar()
plt.xticks([0, 1], ['No Cancer', 'Cancer'])
plt.yticks([0, 1], ['No Cancer', 'Cancer'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/\_classification.py:1344:  
 UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
 0.0 in labels with no predicted samples. Use `zero\_division` parameter to  
 control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/\_classification.py:1344:  
 UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
 0.0 in labels with no predicted samples. Use `zero\_division` parameter to  
 control this behavior.


```
_warn_prf(average, modifier, msg_start, len(result))
```

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/\_classification.py:1344:  
 UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
 0.0 in labels with no predicted samples. Use `zero\_division` parameter to  
 control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

	precision	recall	f1-score	support
1	0.00	0.00	0.00	5
2	0.75	0.78	0.76	49
3	0.43	0.80	0.56	25
4	0.00	0.00	0.00	4

5	1.00	0.93	0.97	15
6	0.86	0.60	0.71	20
7	0.95	0.87	0.90	82
accuracy			0.78	200
macro avg	0.57	0.57	0.56	200
weighted avg	0.79	0.78	0.77	200



```
[ ]: # Install necessary libraries if not already installed
!pip install -q pandas scikit-learn matplotlib

# Import necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report
import matplotlib.pyplot as plt
```

```
[ ]: # Load data (replace 'lung_cancer_data.csv' with your actual data file)
data = pd.read_csv('/content/drive/MyDrive/lung cancer/cancer patient data sets.
↳csv')

# Display first few rows to understand the data
print(data.head())

# Define independent variables (features) and dependent variable (target)
X = data[['Air Pollution', 'Alcohol use']].values # Features
y = data['Genetic Risk'].values # Target variable

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,↳
↳random_state=42)
```

	index	Patient Id	Age	Gender	Air Pollution	Alcohol use	Dust Allergy	\
0	0	P1	33	1	2	4	5	
1	1	P10	17	1	3	1	5	
2	2	P100	35	1	4	5	6	
3	3	P1000	37	1	7	7	7	
4	4	P101	46	1	6	8	7	

	OccuPational Hazards	Genetic Risk	chronic Lung Disease	...	Fatigue	\
0	4	3	2	...	3	
1	3	4	2	...	1	
2	5	5	4	...	8	
3	7	6	7	...	4	
4	7	7	6	...	3	

	Weight Loss	Shortness of Breath	Wheezing	Swallowing Difficulty	\
0	4	2	2	3	
1	3	7	8	6	
2	7	9	2	1	
3	2	3	1	4	
4	2	4	1	4	

	Clubbing of Finger Nails	Frequent Cold	Dry Cough	Snoring	Level
0	1	2	3	4	Low
1	2	1	7	2	Medium
2	4	6	7	2	High
3	5	6	7	5	High
4	2	4	2	3	High

[5 rows x 26 columns]

```
[ ]: # Initialize Random Forest classifier
rfc = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
# Train the model
rfc.fit(X_train, y_train)
```

```
[ ]: RandomForestClassifier(random_state=42)
```

```
[ ]: # Predict on test data
y_pred = rfc.predict(X_test)

# Print classification report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
1	0.00	0.00	0.00	5
2	0.75	0.78	0.76	49
3	0.41	0.52	0.46	25
4	0.00	0.00	0.00	4
5	1.00	0.93	0.97	15
6	0.86	0.60	0.71	20
7	0.88	0.95	0.91	82
accuracy			0.78	200
macro avg	0.55	0.54	0.54	200
weighted avg	0.75	0.78	0.76	200

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```